Microsoft 70-461

Querying Microsoft SQL Server 2012

Version: 20.4

QUESTION NO: 1

You work as a database administrator at TestKing.com. TestKing.com has a SQL Server 2012 database named ProductsDB. The relevant part of the ProductsDB is shown in the following database diagram:

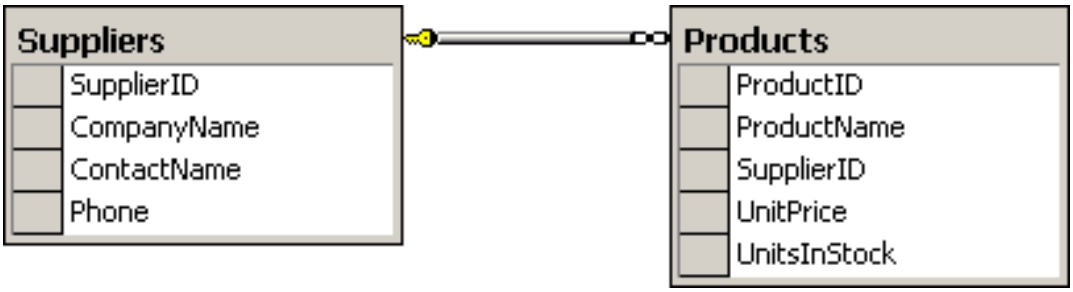

You need to write a Transact-SQL query that display a single row in the following XML format:

<row ProductID="1001" Product="Product Name", Price="24.99" InStock="16" Supplier="Company Name" Contact="Contact Name" Phone="346 959 2215" />

Which of the following SELECT statement would you write?

A.
SELECT ProductID, ProductName AS [Product], UnitPrice AS [Price], UnitsInStock AS [InStock], CompanyName AS [Supplier], ContactName AS [Contact], Phone

FROM Products

INNER JOIN Suppliers ON SupplierID = SupplierID

WHERE ProductID = 1001

FOR XML RAW

B.
SELECT ProductID, ProductName AS [Product], UnitPrice AS [Price], UnitsInStock AS [InStock], CompanyName AS [Supplier], ContactName AS [Contact], Phone

FROM Products

INNER JOIN Suppliers ON SupplierID = SupplierID

WHERE ProductID = 1001

FOR XML

C.
SELECT ProductID, ProductName AS [Product], UnitPrice AS [Price], UnitsInStock AS [InStock], CompanyName AS [Supplier], ContactName AS [Contact], Phone

FROM Products

INNER JOIN Suppliers ON SupplierID = SupplierID

WHERE ProductID = 1001

FOR XML AUTO

D.
SELECT ProductID, ProductName AS [Product], UnitPrice AS [Price], UnitsInStock AS [InStock], CompanyName AS [Supplier], ContactName AS [Contact], Phone

FROM Products

INNER JOIN Suppliers ON SupplierID = SupplierID

WHERE ProductID = 1001

FOR XML AUTO, RAW

Answer: A
Explanation:

QUESTION NO: 2

You work as a database administrator at TestKing.com. TestKing.com has a SQL Server 2012 database named SalesDB. The SalesDB is shown in the following database diagram:

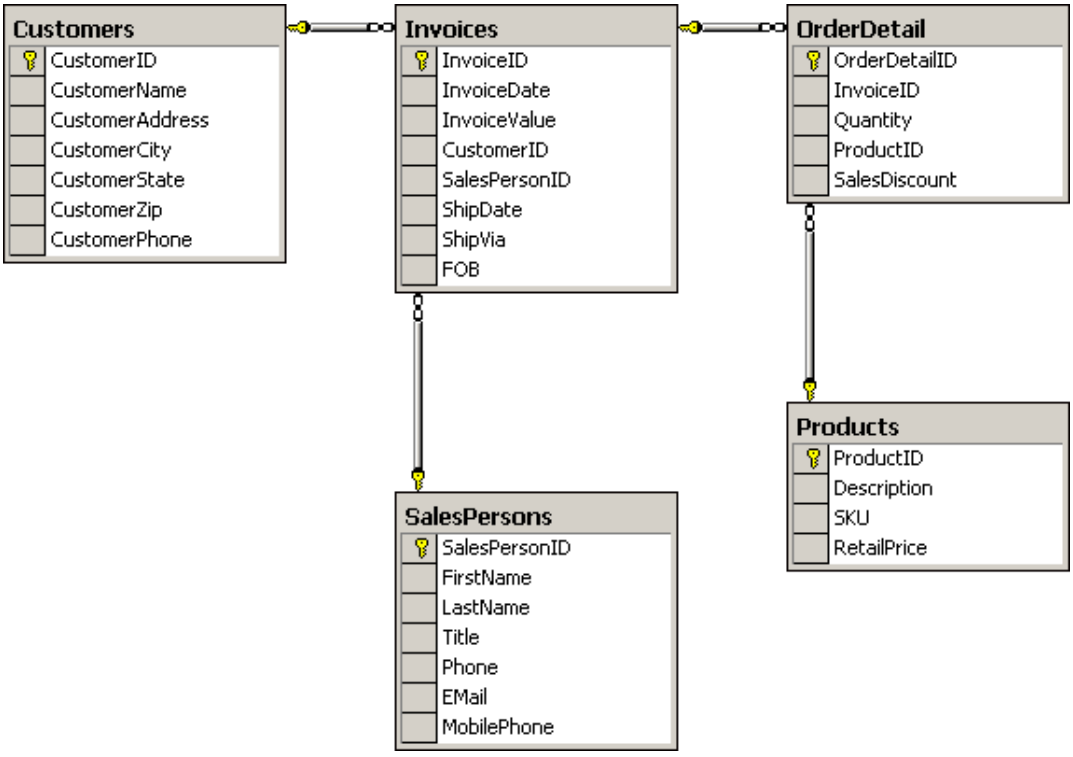

You need to write a Transact-SQL query that display a single row in the following XML format:

<Invoices InvoiceID="1001" Date="2012-10-01T00:00:00", Value="1000.00" Customer="Customer Name" ShippedTo="Customer City" />

Which of the following SELECT statement would you write?

A.
SELECT in.InvoiceID, in.InvoiceDate AS [Date], in.InvoiceValue AS [Value], cu.CustomerName AS [Name], cu.CustomerCity AS [ShippedTo] FROM Invoices AS in

INNER JOIN Customers AS cu ON in.CustomerID = cu.CustomerID

WHERE cu.CustomerID = 1001

FOR XML RAW

B.
SELECT InvoiceID, InvoiceDate AS [Date], InvoiceValue AS [Value], CustomerName AS [Name], CustomerCity AS [ShippedTo] FROM Invoices

INNER JOIN Customers ON Invoices.CustomerID = Customers.CustomerID

WHERE Customers.CustomerID = 1001

FOR XML

C.
SELECT Invoices.InvoiceID, Invoices.InvoiceDate AS [Date], Invoices.InvoiceValue AS [Value], Customers.CustomerName AS [Name], Customers.CustomerCity AS [ShippedTo] FROM Invoices

INNER JOIN Customers ON Invoices.CustomerID = Customers.CustomerID

WHERE Customers.CustomerID = 1001

FOR XML AUTO

D.
SELECT InvoiceID, InvoiceDate AS [Date], InvoiceValue AS [Value], CustomerName AS [Name], CustomerCity AS [ShippedTo] FROM Invoices

INNER JOIN Customers ON Invoices.CustomerID = Customers.CustomerID

WHERE Customers.CustomerID = 1001

FOR XML AUTO, RAW

Answer: A

Explanation:

QUESTION NO: 3

You work as a database administrator at TestKing.com. TestKing.com has a SQL Server 2012 database named ProductsDB. The ProductsDB database is shown in the following database diagram:

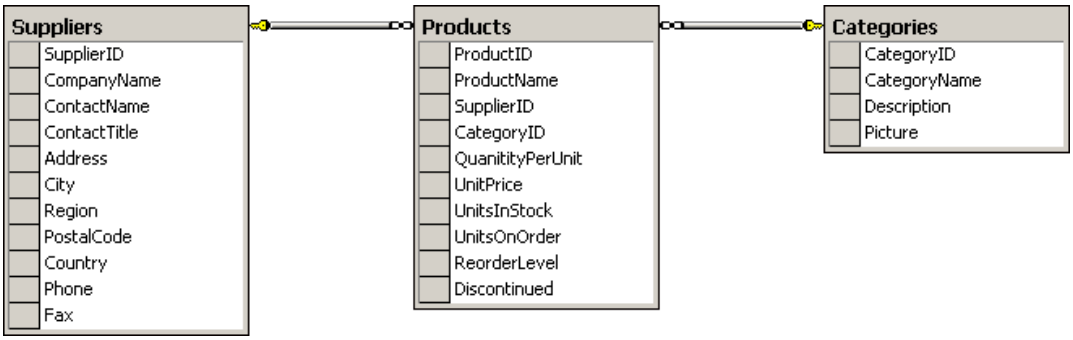

You need to write a Transact-SQL query that displays all the products received by a single supplier in the following XML format:

<Suppliers SupplierID="22" Company="Company Name" ContactNumber="510 250 6400">

<Products ProductID="100" UnitPrice="249.00" UnitsInStock="7" />

<Products ProductID="118" UnitPrice="559.00" UnitsInStock="12" />

</Suppliers>

Which of the following SELECT statement would you write?

A.
SELECT s.SupplierID, s.CompanyName AS [Company], s.ContactNumber, p.ProductID, p.UnitPrice, p.UnitsInStock

FROM Suppliers AS s

INNER JOIN Products AS p ON s.SupplierID = p.SupplierID

WHERE s.SupplierID = 22

FOR XML RAW

B.
SELECT s.SupplierID, s.CompanyName AS [Company], s.ContactNumber, p.ProductID, p.UnitPrice, p.UnitsInStock

FROM Suppliers AS s

INNER JOIN Products AS p ON s.SupplierID = p.SupplierID

WHERE s.SupplierID = 22

FOR XML

C.
SELECT Suppliers.SupplierID, Suppliers.CompanyName AS [Company], Suppliers.ContactNumber, Products.ProductID, Products.UnitPrice, Products.UnitsInStock

FROM Suppliers

INNER JOIN Products ON Suppliers.SupplierID = Products.SupplierID

WHERE Suppliers.SupplierID = 22

FOR XML AUTO

D.
SELECT Suppliers.SupplierID, Suppliers.CompanyName AS [Company], Suppliers.ContactNumber, Products.ProductID, Products.UnitPrice, Products.UnitsInStock

FROM Suppliers

INNER JOIN Products ON Suppliers.SupplierID = Products.SupplierID

WHERE Suppliers.SupplierID = 22

FOR XML AUTO, RAW

Answer: D
Explanation:

QUESTION NO: 4 CORRECT TEXT

You work as a SQL Server 2012 database developer at TestKing.com. You are developing a query for a database driven Web application that allows visitors to vote for the cricket player of the week. The number of votes is stored in a table named WeeklyVotes that has columns named Week, PlayerName, Votes.

You need to write a Transact-SQL query that ranks the top 30 cricket players by the average votes over the last 12 months. You want the top 10 cricket players to have a rank of 1, the next 10 to have a rank of 2, and the last 10 to have a rank of 3.

Which of the following SELECT statement would you use?

To answer, type the correct code in the answer area.

Answer:
SELECT TOP 30 PlayerName,

NTILE (3) OVER (ORDER BY AVG (Votes) DESC) AS AveVotes

FROM WeeklyVotes

GROUP BY PlayerName

QUESTION NO: 5

You work as a database developer at TestKing.com. TestKing.com has a SQL Server 2012 database named SalesDB that has a table named WeeklySales. The WeeklySales table records the sales amount for each of TestKing.com's 20 sales representitives.

You need to write a Transact-SQL query that ranks the sales representatives by the average sales amount for the past year. You want the sales representatives with the same average sales amount to have the same rank with the subsequent rank being skipped.

Which ranking function should you use?

A.
The RANK( ) OVER function.

B.
The NTILE( ) OVER function

C.
The DENSE_RANK( ) OVER function

D.
The ROW_NUMBER( ) OVER function

E.
The FORMAT function.

Answer: A
Explanation:

Ref: http://msdn.microsoft.com/en-us/library/ms189798.aspx

QUESTION NO: 6

You work as a SQL Server 2012 database developer at TestKing.com. You are developing a query for a database driven Web application that allows visitors to vote for the cricket player of the week. The number of votes is stored in a table named WeeklyVotes that has columns named Week, PlayerName, Votes.

You need to write a Transact-SQL query that returns the cricket player that received the most votes for each week, as well as the number of votes they received.

Which of the following SELECT statement would accomplish this task?

A.
SELECT PlayerName, Votes

FROM (SELECT PlayerName, Votes,

RANK () OVER (PARTITION BY PlayerName ORDER BY Votes ASC) AS Rank

FROM WeeklyVotes) AS tmp

WHERE Rank = 1

B.
SELECT PlayerName, Votes

FROM (SELECT PlayerName, Votes,

RANK() OVER (PARTITION BY Week ORDER BY Votes DESC) AS Rank FROM WeeklyVotes) AS tmp

WHERE Rank = 1

C.
SELECT PlayerName, Votes

FROM (SELECT TOP 1 PlayerName, Votes,

RANK () OVER (PARTITION BY PlayerName ORDER BY Votes ASC) AS Rank

FROM WeeklyVotes

ORDER BY Rank) AS tmp

D.
SELECT PlayerName, Votes

FROM (SELECT TOP 1 PlayerName, Votes,

RANXO OVER (PARTITION BY PlayerName ORDER BY Votes DESC) AS Rank

FROM WeeklyVotes

ORDER BY Rank) AS tmp

Answer: B
Explanation:

QUESTION NO: 7

You work as a database developer at TestKing.com. TestKing.com has a SQL Server 2012 database named SalesDB that has a table named Inventory.

The Inventory table has three columns named ProductID, InStore and InWarehouse. The ProductID column is the primary key and is linked to the Products table. The InStore column stores the quantity of a product that is held at TestKing.com's retail shop, while the InWarehouse column stores the quantity of a product that is held at TestKing.com's warehouse.

You need to add a computed column that stores the total number of a product that TestKing.com has.

What Transact-SQL statements would accomplish this task?

A.
ALTER TABLE Inventory

ADD TotalProducts AS (InStore + InWarehouse) PERSISTED

B.
ALTER TABLE Inventory

ADD TotalProducts int SPARSE NOT NULL

C.
ALTER TABLE Inventory

ADD TotalProducts AS SUM (ALL) OVER (GROUP BY InStore, InWarehouse) PERSISTED

D.
DROP TABLE Inventory

GO

CREATE TABLE Inventory

(

ProductID int NOT NULL PRIMARY KEY,

InStore int NOT NULL,

InWarehouse int NOT NULL,

TotalProducts AS SUM (InStore, InWarehouse) PERSISTED

)

Answer: A
Explanation:

Ref: http://www.kodyaz.com/articles/sql-server-computed-column-calculated-column-sample.aspx

QUESTION NO: 8

TestKing.com has a SQL Server 2012 database instance that hosts a database named ComDB. The ComDB database has a table named Partners that was created using the following Transact-SQL code:

CREATE TABLE [dbo].[Partners]

(

[CompanyID] [int] NOT NULL,

[CompanyName] [nvarchar] (50) NOT NULL,

[Location] [nvarchar] (50) NOT NULL,

[ContactName] [nvarchar] (50) NOT NULL,

[Email] [nvarchar] (50) NOT NULL,

[Phone] [nvarchar] (10) NOT NULL,

CONSTRAINT [PK_Partners] PRIMARY KEY CLUSTERED

(

[CompanyID] ASC

)

ON PRIMARY

)

You want to create a FOR UPDATE trigger that will track changes to the ContactName and Phone columns.

Which of the following statements should you use in the trigger definition?

A.
IF COLUMNS_UPDATED (ContactName, Phone)

B.
IF COLUMNS_UPDATED (ContactName) OR COLUMNS_UPDATED (Phone)

C.
IF UPDATED (ContactName, Phone).

D.
IF UPDATED (ContactName) OR UPDATED (Phone)

Answer: D
Explanation:

QUESTION NO: 9

You are the database administrator of a SQL Server 2012 database infrastructure at TestKing.com.

You need to optimize a very large database table that contains several million rows of data by designing a view based on the table. The view must allow users to perform aggregations on several columns.

How should you configure the view to ensure optimal performance?

A.
You should create the view as an indexed view.

B.
You should create a clustered index on the view.

C.
You should make use of a stored procedure to return that data.

D.

You should make use of a table-valued function.

Answer: A
Explanation:

QUESTION NO: 10

You are the database developer at TestKing.com. TestKing.com has a SQL Server 2012 database infrastructure that has a database named ComDB with a table named Partners.

The Partners table was created using the following Transact-SQL code:

```
CREATE TABLE [dbo].[Partners]
(
[CompanyID] [int] NOT NULL PRIMARY KEY,
[CompanyName] [varchar] (150) NOT NULL,
[Location] [varchar] (150) NOT NULL,
[ContactName] [varchar] (150) NOT NULL,
[Email] [varchar] (150) NOT NULL,
[Phone] [varchar] (10) NOT NULL
)
```

You develop a new table named Events using the following Transact-SQL code:

```
CREATE TABLE [dbo].[Events]
(
[EventID] [int] NOT NULL PRIMARY KEY,
[CompanyID] [int] NOT NULL,
[EventDescription] [varchar] (2500),
[EventCordinator] [varchar] (150) NOT NULL
)
```

How would you guarantee that values in the Events.CompanyID column already exist in the Partners.CompanyID column?

A.
You should add a Foreign Key Constraint on the Events table.

B.
You should add a Check Constraint on the Events table.

C.
You should add a Unique Constraint on the Events table.

D.
You should specify Events.CompanyID as a spars column.

E.
You should change the Events table to a partitioned table.

Answer: A
Explanation:

Ref: http://msdn.microsoft.com/en-us/library/ms179610.aspx

QUESTION NO: 11

TestKing.com has a SQL Server 2012 database infrastructure that has a database named ComDB.

You have created a view using the following Transact-SQL code:

```
CREATE VIEW tkCommunications

(Type, CompanyID, CompanyName, Location, ContactName, Email, Phone)

AS

SELECT 'Clients' AS Type, CompanyID, CompanyName, Location, ContactName, Email, Phone

FROM CommList

WHERE Relation = 'Client'

SELECT 'Partners' AS Type, CompanyID, CompanyName, Location, ContactName, Email, Phone
```

```
FROM CommList

WHERE Relation = 'Partner'

SELECT 'Guests' AS Type, CompanyID, CompanyName, Location, ContactName, Email, Phone

FROM CommList

WHERE Relation = 'Guests'

GO
```

You want the view to be used to edit all columns except the CompanyID, CompanyName and Location columns.

What should you implement on the view?

A.
You should consider implementing an AFTER UPDATE trigger.

B.
You should consider implementing an Index.

C.
You should consider implementing an INSTEAD OF UPDATE trigger.

D.
You should consider implementing a CHECK constraint.

Answer: C
Explanation:

QUESTION NO: 12

You work as a database administrator at TestKing.com. TestKing.com has a SQL Server 2012 database named SalesDB. The SalesDB database is shown in the following database diagram:

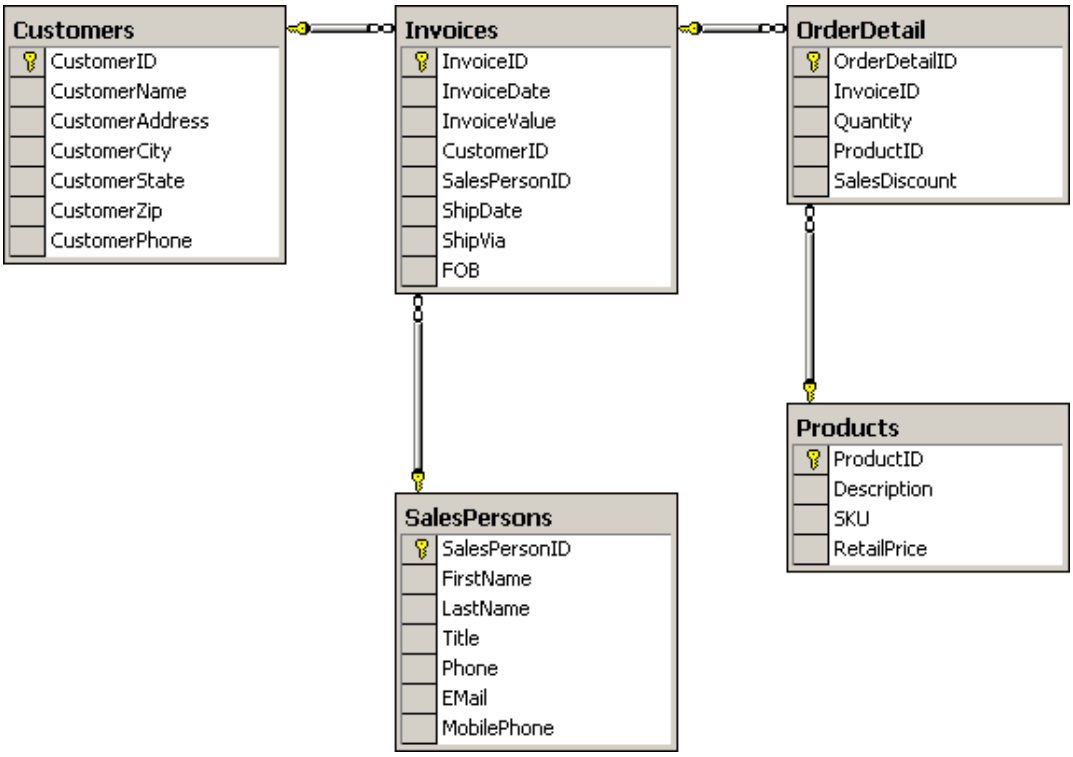

You create a view on the SalesDB using the following Transact-SQL code:

CREATE VIEW SalesV

WITH SCHEMABINDINGS

AS

SELECT Products.ProductID, Invoices.InvoiceDate, SUM (Products.RetailPrice * OrderDetail.Quantity * OrderDetail.SalesDiscount) AS Price

FROM OrderDetail INNER JOIN Products ON

OrderDetail.ProductID = Products.ProducID

INNER JOIN Invoices ON

OrderDetail.InvoiceID = Invoices.InvoiceID

GROUP BY Products.ProductID, Invoices.InvoiceDate

GO

How should you alter this view to allow users to update data through the SalesV?

A.
You should add a CHECK constraint to the SalesV view.

B.
You should add an INSTEAD OF trigger to the SalesV view.

C.
You should add a clustered index to the SalesV view.

D.
You should add an AFTER UPDATE trigger to the SalesV view.

E.
Create a columnstore index on all columns used in the SalesV view.

Answer: B
Explanation:

QUESTION NO: 13 CORRECT TEXT

You are employed as a SQL Server 2012 database developer at TestKing.com. TestKing.com has a SalesDB database with a view named SalesV. The SalesV view was created using the following Transact-SQL code:

```
CREATE VIEW SalesDB.ProductsSalesV

AS

SELECT OrderID, ProductID, ShipDate, OrderDate, Amount

FROM SalesDB.Orders;
```

You want to create an inline table-valued function named fn_testking that accepts a @ProductID parameter of the integer data type. The inline table-valued function should also allow for sales orders for each product to be listed by the latest sale.

How would you create this inline table-valued function?

To answer, type the correct code in the answer area.

Answer:

```
CREATE FUNCTION SalesDB.fn_testking ( @ProductID int )

RETURNS TABLE

AS

RETURN

(
```

```
SELECT OrderID, ProductID, ShipDate, OrderDate, Amount

FROM Sales. ProductsSalesV

WHERE ProductID = @ProductID

ORDER BY OrderDate DESC

);
```

QUESTION NO: 14 CORRECT TEXT

You are employed as a SQL Server 2012 database developer at TestKing.com. TestKing.com has a database named SalesDB with tables named Customer and Orders. The Customer and Orders tables were created using the following Transact-SQL code:

```
CREATE TABLE SalesDB.Customers

(

CustomerID int NOT NULL PRIMARY KEY,

FirstName varchar (150) NOT NULL,

LastName varchar (150) NOT NULL,

Address1 varchar (200) NOT NULL,

Address2 varchar (200) NULL,

City varchar (100) NOT NULL,

StateCode varchar (2) NOT NULL,

Zip varchar (5) NOT NULL,

Phone varchar (10) NOT NULL

)

GO

CREATE TABLE SalesDB.Orders

(

OrderID int NOT NULL PRIMARY KEY,

CustomerID int NOT NULL,
```

OrderDate datetime NOT NULL,

ShipDate datetime NOT NULL,

CustomerID int NOT NULL,

SalesRepID int NOT NULL

)

GO

You must now create an OrderDetails table as shown in the following database diagram:

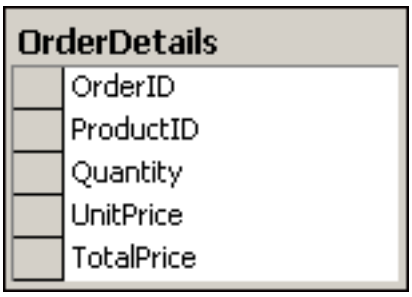

The TotalPrice column must be a computed column based on the product of the UnitPrice and Quantity columns and the data must be stored in the table.

How would you create this table?

To answer, type the correct code in the answer area.

Answer:
CREATE TABLE SalesDB.OrderDetails

(

OrderID int NOT NULL,

ProductID int NOT NULL,

Quantity int NOT NULL,

UnitPrice money NOT NULL,

TotalPrice AS (Quantity * UnitPrice) PERSISTED

)

QUESTION NO: 15

You work as a SQL Server 2012 database developer at TestKing.com. TestKing.com has a database named SalesDB.

You are developing a stored procedure that takes a parameter named @date that uses the varchar datatype. The @date parameter must be compared to the value in a datetime column named OrderDate.

Which of the following WHERE clauses would be the most efficient WHERE clause to use?

A.
WHERE OrderDate = CAST(datetime,@date)

B.
WHERE OrderDate = CONVERT(datetime,@date)

C.
WHERE OrderDate =@date

D.
WHERE OrderDate = CAST(@date AS datetime)

E.
WHERE OrderDate = PARSE(@date AS Date)

Answer: C
Explanation:

QUESTION NO: 16

You work as a database developer at TestKing.com. You want to create a Transact-SQL query will call a table-valued function for every row the query returns.

How would you accomplish this task?

A.
You should make use of a UNION.

B.
You should make use of a CONVERT function.

C.
You should make use of an INNER JOIN.

D.
You should make use of a Trigger.

E.
You should make use of a CAST function.

F.
You should make use of an OUTER JOIN.

G.
You should make use of a CROSS APPLY.

H.
You should make use of the FORMAT function.

Answer: G
Explanation:

QUESTION NO: 17

Which of the following datatypes has a fixed precision and a scale of six digits?

A.
Double

B.
Money

C.
Int

D.
Numeric

E.
SmallInt

F.
VarInt

G.
Float

Answer: D

Explanation:

QUESTION NO: 18

You work as a database developer at TestKing.com. TestKing.com has a SQL Server 2012 database named SalesDB as illustrated in the following database diagram:

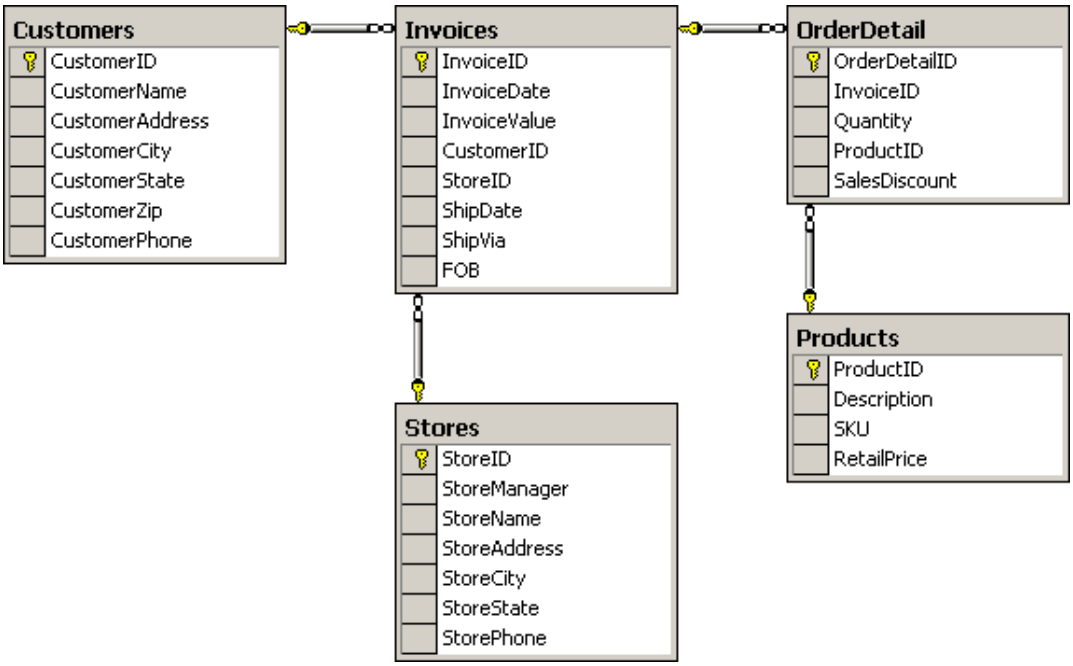

TestKing.com has retail stores in a few major cities across the country. The company wants a list of Customers who live in a city that does not have a TestKing.com store, along with the customer's City. The result set must be sorted alphabetically by City name.

Which of the following Transact-SQL statements would return the required information?

A.
SELECT CustomerName, CustomerCity

FROM Customers

WHERE CustomerCity NOT EXISTS (SELECT StoreCity FROM Stores)

ORDER BY CustomerCity

B.
SELECT CustomerName, CustomerCity

FROM Customers

WHERE CustomerCity < > ALL (SELECT StoreCity FROM Stores)

ORDER BY StoreCity

C.
SELECT CustomerName, CustomerCity

FROM Customers

WHERE CustomerCity < > ANY (SELECT StoreCity FROM Stores)

ORDER BY CustomerCity

D.
SELECT CustomerName, CustomerCity

FROM Customers

WHERE CustomerCity NOT IN (SELECT StoreCity FROM Stores)

ORDER BY StoreCity

Answer: C
Explanation:

QUESTION NO: 19

You work as a database developer at TestKing.com. TestKing.com has a SQL Server 2012 database infrastructure with several databases. You have permissions on two of these databases, namely, CommDB and SalesDB.

You need to develop a stored procedure named dbo.tk_addEntry in a database named AssetsDB. The dbo.tk_addEntry stored procedure will run stored procedures in CommDB and SalesDB.

How should you design the stored procedure so that callers that do not paver permissions on CommDB and SalesDB can run the dbo.tk_addEntry stored procedure successfully?

A.
You should configure the stored procedure to call the xp_cmdshell extended stored procedure.

B.
You should configure the stored procedure to call the sp_configure system stored procedure.

C.
You should assign users permission to the stored procedure.

D.
You should include the EXECUTE AS CALLER clause when creating the stored procedure.

E.
You should include the EXECUTE AS OWNER clause when creating the stored procedure.

Answer: E
Explanation:

QUESTION NO: 20 CORRECT TEXT

You are the database developer at TestKing.com. TestKing.com has a SQL Server 2012 database infrastructure that has a database named ComDB with tables named Partners and Events. These tables were created using the following Transact-SQL code:

```
CREATE TABLE [dbo].[Partners]
(
[CompanyID] [int] NOT NULL PRIMARY KEY,
[CompanyName] [varchar] (150) NOT NULL,
[Location] [varchar] (150) NOT NULL,
[ContactName] [varchar] (150) NOT NULL,
[Email] [varchar] (150) NOT NULL,
[Phone] [varchar] (10) NOT NULL
)

CREATE TABLE [dbo].[Events]
(
[EventID] [int] NOT NULL PRIMARY KEY,
[CompanyID] [int] NOT NULL,
[EventDescription] [nvarchar] (MAX),
[EventDate] [nvarchar] (50) NOT NULL,
[EventCordinator] [nvarchar] (150) NOT NULL
)
```

You add a foreign key relationship between the two tables on the CompanyID column.

You need to develop a stored procedure named sp_coEvents that retrieves CompanyName for all partners and the EventDate for all events that they have coordinated.

To answer, type the correct code in the answer area.

Answer:
A

CREATE PROCEDURE sp_coEvent

AS

SELECT CompanyName, EventDate

FROM Partners JOIN Events ON

Partners.CompanyID = Events.CompanyID

QUESTION NO: 21

You work as a database developer at TestKing.com. TestKing.com has a SQL Server 2012 database infrastructure with a very large database named SalesDB. You create a new table named SalesHistory that will hold historical data from the SalesDB database.

You need to perform a batch update from the SalesDB database to the SalesHistory table. You want the status information from the batch process to be logged to a SQL Server table that must be created by the batch process.

How would you accomplish this task?

A.
You should make use of the FORMAT function.

B.
You should make use of the CONVERT function.

C.
You should make use of a scalar user-defined function.

D.
You should make use of an inline function.

E.
You should make use of a table-valued function.

F.
You should make use of a stored procedure.

Answer: C
Explanation:

QUESTION NO: 22

You are developing a SQL Server 2012 database for TestKing.com. You need to create a computed column that returns the data by referencing another table using an INNER JOIN.

How would you accomplish this?

A.
You should make use of the FORMAT function.

B.
You should make use of a scalar user-defined function.

C.
You should make use of an inline function.

D.
You should make use of a table-valued user-defined function.

E.
You should make use of a stored procedure.

Answer: C
Explanation:

QUESTION NO: 23

Which of the following statements regarding SQL Server 2012 objects is TRUE?

A.
A user-defined data type can accept an input variable and return a table of results but cannot be

used within a view.

B.
A scalar function can accept an input variable and return a table of results but cannot be used within a view.

C.
A table-valued function can accept an input variable and return a table of results but cannot be used within a view.

D.
A table-valued type can accept an input variable and return a table of results but cannot be used within a view.

Answer: C
Explanation:

QUESTION NO: 24

You work as a SQL Server 2012 database developer at TestKing.com. TestKig.com has a database named DataDB.

You are developing a complex stored procedure named sp_updater that will use a single transaction to update several tables in the DataDB database.

You are concerned about data integrity and incomplete updates should the sp_updater stored procedure cause a run-time error.

To mitigate this potential problem you want the transaction to terminate and the transaction to be rolled back if the sp_updater stored procedure raises a run-time error.

How would you accomplish this task?

A.
You should make use of the SET XACT_ABORT ON statement in the stored procedure.

B.
You should have the stored procedure run in the SERIALIZABLE ISOLATION LEVEL.

C.

You should make use of a LOOP hint in the stored procedure.

D.
You should have the stored procedure run in the SNAPSHOT ISOLATION LEVEL.

E.
You should make use of an INSTEAD OF UPDATE trigger in the stored procedure.

Answer: A
Explanation:

Ref: http://msdn.microsoft.com/en-us/library/ms188792

QUESTION NO: 25

You work as a SQL Server 2012 database developer at TestKing.com.

Testking.com has a database SalesDB with a large Orders table. You create a heap named OldData that will store historical data from the Orders table.

You need to write a Transact-SQL query that will insert rows of data from the Orders table that are marked as closed and are more than six months old.

Which of the following table hints should you use in your query if you want to optimize transaction logging and locking for the query?

A.
You should make use of the READPAST hint.

B.
You should make use of the HOLDLOCK hint.

C.
You should make use of the READCOMMITTED hint.

D.
You should make use of the NOLOCK hint.

E.
You should make use of the TABLOCK hint.

F.
You should make use of the UPDLOCK hint.

Answer: E

Explanation:

QUESTION NO: 26

You are employed as a SQL Server 2012 database developer at TestKing.com. You have a stored procedure that is executed quite often. The stored procedure joins data from two tables.

TestKing.com users report that the stored procedure takes a long time to execute. You analyze the query plan and find that the stored procedure often makes use of table scans rather than indexes when the estimated rows do not match the actual rows on one of the tables.

How would you optimize the performance of the stored procedure?

A.
You should make use of the KEEPIDENTITY table hint in the stored procedure.

B.
You should make use of the KEEPDEFAULTS table hint in the stored procedure.

C.
You should make use of the IGNORE_CONSTRAINTS table hint in the stored procedure.

D.
You should make use of the FORCESEEK table hint in the stored procedure.

E.
You should update statistics on the tables queried by the stored procedure.

Answer: D

Explanation:

QUESTION NO: 27

You work as a SQL Server 2012 database developer at TestKing.com. You are developing a database driven Web application. The application executes a store procedure based on the location of the web user.

The location of the Web user is determined by IP Geolocation. You want to develop a process that

will execute the correct stored procedure for every Web user based on the user's location.

How would you accomplish this?

A.
You should make use of a foreach SQLCLR statement.

B.
You should make use of a scalar user-defined function.

C.
You should make use of an inline function.

D.
You should make use of a cursor.

Answer: D
Explanation:

QUESTION NO: 28

You work as a SQL Server 2012 database administrator at TestKing.com. You have developing a database named WebAnDB that will be used by a web site analysis application. The WebAnDB database has a table named Visitors that stores date and time data in a column named Accessed.

You must now develop a stored procedure that will insert data into the Accessed column. You want the stored procedure to store time zone data as well.

How would you accomplish this?

A.
You should make use of a scalar user-defined function.

B.
You should make use of the SET CONTEXT_INFO statement in the stored procedure.

C.
You should make use of the DATETIMEOFFSET data type.

D.
You should make use of the TODATETIMEOFFSET function.

E.

You should make use of the SET FORCEPLAN ON statement in the stored procedure.

F.
You should make use of a cursor.

Answer: D
Explanation:

QUESTION NO: 29

You work as a SQL Server 2012 database developer at TestKing.com.

You are developing a database that will be used by a Web application. The database will store small multimedia files in several tables. The largest multimedia file is 975 kB. These multimedia files will be retrieved quite often.

How would you store these files?

A.
You should make use of a cursor.

B.
You should make use of a row-level compression.

C.
You should make use of the FileStream data type.

D.
You should make use of the VARBINARY data type.

Answer: D
Explanation:

QUESTION NO: 30

You work as a database administrator at TestKing.com. You are developing a database that will be used by a web site analysis application name TKWeb1.

The TKWeb1 application must display the date and time each visitor visits a page on a website as well as the date and time they leave that web page. This data needs to be displayed in different date and time formats.

How would you accomplish this?

A.
You should make use of a scalar user-defined function.

B.
You should make use of the SET CONTEXT_INFO statement in the stored procedure.

C.
You should make use of the DATETIMEOFFSET data type.

D.
You should make use of the FORMAT function.

E.
You should make use of the SET FORCEPLAN ON statement in the stored procedure.

F.
You should make use of a cursor.

Answer: D
Explanation:

QUESTION NO: 31

You work as a database developer at TestKing.com. TestKing.com has a SQL Server 2012 database named SalesDB that has a table named Inventory.

The Inventory table has three columns named ProductID, InStore and InWarehouse. The ProductID column is the primary key and is linked to the Products table. The InStore column stores the quantity of a product that is held at TestKing.com's retail shop, while the InWarehouse column stores the quantity of a product that is held at TestKing.com's warehouse.

You need to add a computed column that is the sum of values in the InStore and InWarehoue columns for each product.

What Transact SQL statement would accomplish this task?

A.
ALTER TABLE Inventory

ADD TotalProducts AS (InStore + InWarehouse)

B.
ALTER TABLE Inventory

ADD TotalProducts int SPARSE NOT NULL

C.
ALTER TABLE Inventory

ADD TotalProducts AS SUM (ALL) OVER (GROUP BY InStore, InWarehouse)

D.
DROP TABLE Inventory

GO

CREATE TABLE Inventory

(

ProductID int NOT NULL PRIMARY KEY,

InStore int NOT NULL,

InWarehouse int NOT NULL,

TotalProducts AS SUM (InStore, InWarehouse)

)

Answer: A
Explanation:

Ref: http://www.kodyaz.com/articles/sql-server-computed-column-calculated-column-sample.aspx

QUESTION NO: 32

You work as a database administrator at TestKing.com. TestKing.com has a SQL Server 2012 database named SalesDB. The SalesDB database is shown in the following database diagram:

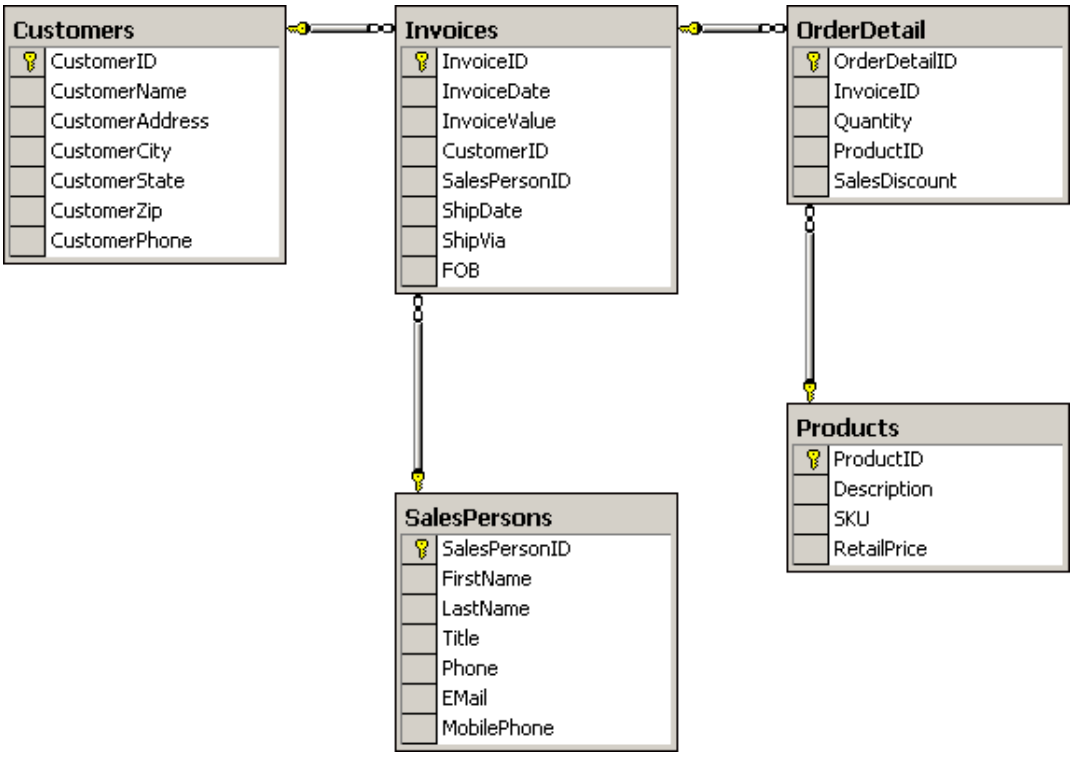

You create a view on the SalesDB database using the following Transact-SQL code:

CREATE VIEW SalesV

WITH SCHEMABINDINGS

AS

SELECT Products.ProductID, Invoices.InvoiceDate, SUM (Products.RetailPrice * OrderDetail.Quantity * OrderDetail.SalesDiscount) AS Price

FROM OrderDetail INNER JOIN Products ON

OrderDetail.ProductID = Products.ProducID

INNER JOIN Invoices ON

OrderDetail.InvoiceID = Invoices.InvoiceID

GROUP BY Products.ProductID, Invoices.InvoiceDate

GO

You want the SalesV view to persist data to disk in order to improve performance.

How would you accomplish this task?

A.
You should add a clustered index to the SalesV view.

B.
You should create a columnstore index on all columns used in the SalesV view.

C.
You should drop and recreate the SalesV view as a system view.

D.
You should drop and recreate the SalesV view as a partitioned view.

Answer: C
Explanation:

Ref: http://msdn.microsoft.com/en-us/library/ms190174

QUESTION NO: 33

You work as a database developer at TestKing.com. You need to store Microsoft Word documents in a SQL Server 2012 database table.

The Word documents must not be accessed by Windows applications that do not use Transact-SQL queries.

How would accomplish this task?

A.
You should store the documents using the VARBINARY(MAX) datatype.

B.
You should store the documents using the VARCHAR(MAX) datatype.

C.
You should store the documents using the FILESTREAM datatype.

D.
You should store the documents using the NVARCHAR(MAX) datatype.

Answer: A
Explanation:

QUESTION NO: 34

You work as a database developer at TestKing.com. TestKing.com has a SQL Server 2012 database named SalesDB with a table named Invoices.

Application developers are developing several in-house applications that will access the Invoices table. You need to develop a solution that will allow the applications to access the table indirectly while still allowing them to update the Invoice table.

How would you accomplish this task?

A.
You should create a view on the Invoices table.

B.
You should create a columnstore index on all columns used by the applications.

C.
You should allow the applications access to the Invoices table via stored procedures.

D.
You should drop and recreate the Invoices table as a partitioned table.

Answer: C
Explanation:

QUESTION NO: 35

Which of the following can be used to protect the code in a stored procedure?

A.
The ENCRYPBYKEY statement.

B.
The ENCRYPBYASYMKEY statement.

C.
The SET TRUSTWORTHY ON option.

D.
The SET XACT_ABORT ON statement.

E.
The ENCRYPTBYPASSPHRASE statement.

F.
The ENCRYPTBYCERT statement.

G.
The SIGNBYASYMKEY statement.

H.
The CRYPT_GEN_RANDOM statement.

Answer: A
Explanation:

QUESTION NO: 36

You work as a database administrator at TestKing.com. TestKing.com has a SQL Server 2012 database named SalesDB. The SalesDB database is shown in the following database diagram:

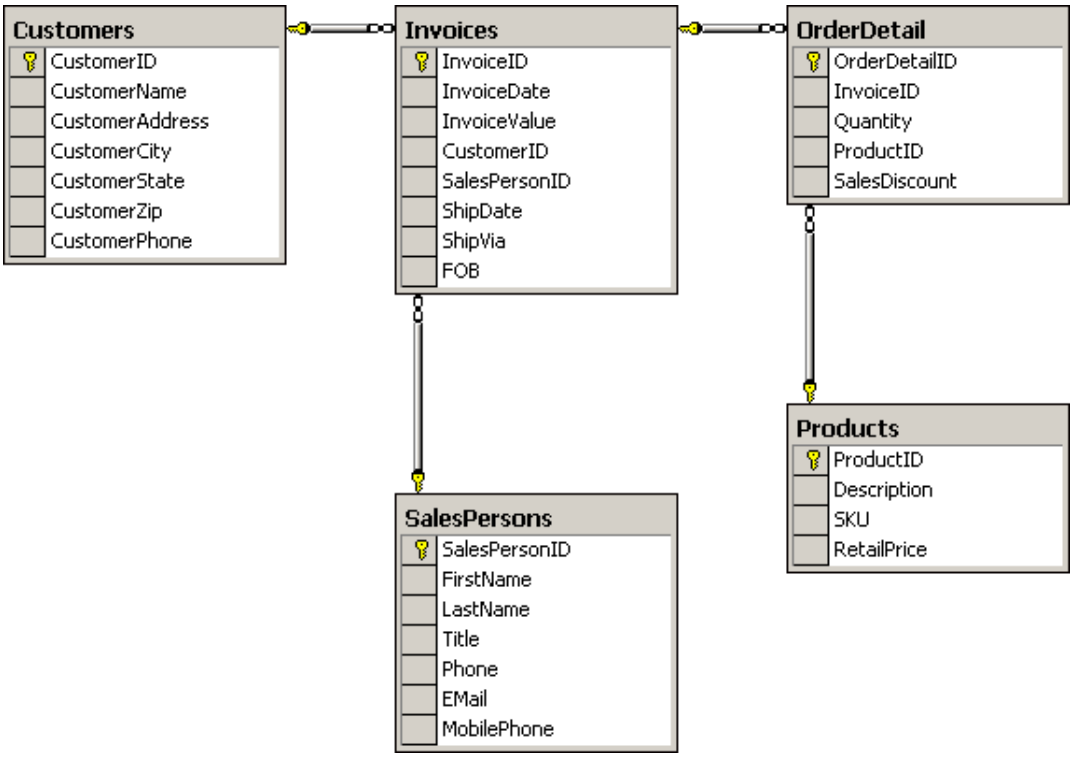

You need to write a Transact-SQL query that displays all Customers, whether they have invoices or not. The query must also display the InvoiceDate for a Customer that has an Invoice.

How would you accomplish this task?

A.
You should make use of a UNION.

B.
You should make use of an INNER JOIN.

C.

You should make use of a CROSS JOIN.

D.
You should make use of an OUTER JOIN.

E.
You should make use of a CROSS APPLY.

Answer: D
Explanation:

QUESTION NO: 37 CORRECT TEXT

You are employed as a SQL Server 2012 database developer at TestKing.com. TestKing.com has a database named HRDB with tables named Staff and SalaryHistory. The Staff and SalaryHistory tables were created using the following Transact-SQL code:

```
CREATE TABLE HRDB.Staff
(
StaffID int NOT NULL PRIMARY KEY,
FirstName nvarchar(50) NOT NULL,
LastName nvarchar(50) NOT NULL,
JobTitle nvarchar(50) NOT NULL,
StaffAddress nvarchar (200) NULL,
StaffCity nvarchar(50) NOT NULL,
StaffState nvarchar(50) NOT NULL,
StaffPhone varchar (10) NOT NULL,
Salary smallmoney NOT NULL
)
GO
CREATE TABLE HRDB.SalaryHistory
(
```

StaffID int NOT NULL,

IncreaseDate date DEFAULT GETDATE(),

PrevSalary smallmoney NOT NULL,

NewSalary smallmoney NOT NULL

)

GO

You must write a Transact-SQL query to affect a salary increase of 6.5% for all staff members with a JobTitle of Support Technician. Your query must also update the data in the SalaryHistory table.

What Transact SQL statement would accomplish this task?

To answer, type the correct code in the answer area.

Answer:
UPDATE Staff

SET Salary = Salary * 1.065

WHERE JobTitle = 'Support Technician'

OUTPUT inserted.StaffID, deleted.Salary, inserted.Salary

INTO SalaryHistory(StaffID, PrevSalary, NewSalary)

QUESTION NO: 38

You work as a SQL Server 2012 database developer at TestKing.com. You are developing a batch process that accepts an input parameter and allows the result set to be joined with a table.

How would you accomplish this?

A.
You should make use of a stored procedure to perform the batch process.

B.
You should make use of a table-valued user-defined function.

C.
You should make use of the SET XACT_ABORT ON statement.

D.
You should make use of a scalar user-defined function.

E.
You should make use of a schema-bound user-defined function.

Answer: B
Explanation:

Ref: h http://msdn.microsoft.com/en-us/library/ms191007.aspx

QUESTION NO: 39

You work as a database administrator at manufacturing company named TestKing.com. TestKing.com has a SQL Server 2012 database named ProductionDB. The ProductionDB database has a table named Sites that was created using the following Transact-SQL code:

CREATE TABLE Sites (

SiteID int NOT NULL PRIMARY KEY,

Location int NOT NULL,

Manager nvarchar(200) NOT NULL,

Week smallint NOT NULL,

ItemsProduced int NOT NULL)

You want to write the Transact-SQL query that returns the number of items produced at each location for each week. In addition, you want the result set to include a column named PrevItemsProduced that holds the number of items produced at each location for the week before.

What Transact SQL statement would accomplish this task?

A.
SELECT Location, Week, ItemsProduced,

LEAD(ItemsProduced, 1, 0) OVER (PARTITION BY Location ORDER BY Week) AS PrevItemsProduced

FROM Sites

B.
SELECT Location, Week, ItemsProduced,

FIRST_VALUE(ItemsProduced) OVER (PARTITION BY Location ORDER BY Week) AS PrevItemsProduced

FROM Sites

C.
SELECT Location, Week, ItemsProduced,

LAG(ItemsProduced, 1, 0) OVER (PARTITION BY Location ORDER BY Week) AS PrevItemsProduced

FROM Sites

D.
SELECT Location, Week, ItemsProduced,

LAST_VALUE(ItemsProduced) OVER (PARTITION BY Location ORDER BY Week) AS PrevItemsProduced

FROM Sites

E.
SELECT Location, Week, ItemsProduced,

CUME_DIST( ) OVER (PARTITION BY Location ORDER BY Week) AS PrevItemsProduced

FROM Sites

Answer: C
Explanation:

QUESTION NO: 40

You work as a SQL Server 2012 database developer at TestKing.com. You are developing a stored procedure for TestKing.com's e-Commerce application.

Your stored procedure must display unique values from one column in multiple columns in the result set.

How would you accomplish this?

A.
You should make use of the OUTER APPLY operator.

B.
You should make use of a dynamic cursor.

C.
You should make use of the PIVOT operator.

D.
You should make use of the CROSS APPLY operator.

E.
You should make use of the UNPIVOT operator.

F.
You should make use of a keyset cursor.

Answer: C
Explanation:

QUESTION NO: 41

You work as a database developer at TestKing.com. TestKing.com has a SQL Server 2012 database named SalesDB as illustrated in the following database diagram:

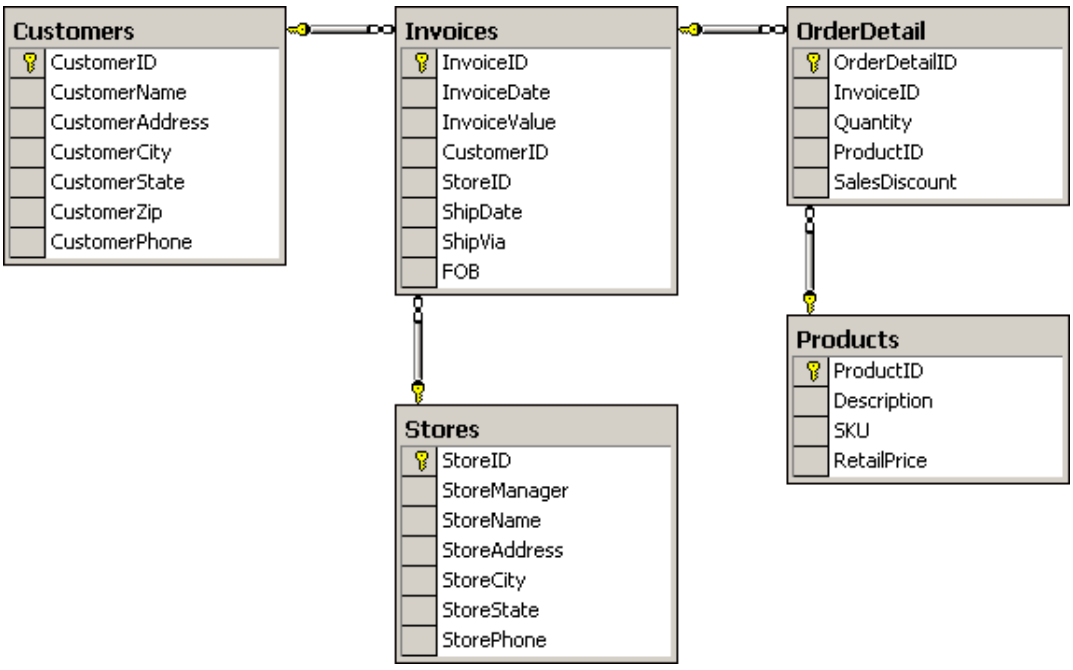

TestKing.com has retail stores in a few major cities across the country. The company wants to ascertain whether it would be advantageous to open a store in other cities based on feedback from its customers.

You are required to provide the company's CEO with a list of Customers who live in a city that does not have a TestKing.com store, along with the customer's Phone Number and the customer's City, and arranged alphabetically by City name.

Which of the following Transact-SQL statements would return the required information?

A.
SELECT CustomerName, CustomerCity, CustomerPhone

FROM Customers

WHERE CustomerCity NOT EXISTS (SELECT StoreCity FROM Stores)

ORDER BY CustomerCity

B.
SELECT CustomerName, CustomerCity, CustomerPhone

FROM Customers

WHERE CustomerCity < > ALL (SELECT StoreCity FROM Stores)

ORDER BY StoreCity

C.
SELECT CustomerName, CustomerCity, CustomerPhone

FROM Customers

WHERE CustomerCity < > ANY (SELECT StoreCity FROM Stores)

ORDER BY CustomerCity

D.
SELECT CustomerName, CustomerCity, CustomerPhone

FROM Customers

WHERE CustomerCity NOT IN (SELECT StoreCity FROM Stores)

ORDER BY StoreCity

Answer: C
Explanation:

QUESTION NO: 42

You work as a database developer at TestKing.com. You are developing a SQL server 2012 database for TestKing.com's e-Commerce application. The application allows TestKing.com employees from different regions to execute a store procedure based on their location.

The location of the TestKing.com employees are determined by an input parameter named @location. You want to develop a process that will execute the correct stored procedure for every TestKing.com employee based on their location.

How would you accomplish this?

A.
You should make use of a client cursor.

B.
You should make use of a static cursor.

C.
You should make use of a forward-only cursor.

D.
You should make use of a dynamic cursor.

E.
You should make use of a keyset cursor.

Answer: E
Explanation:

Ref: http://msdn.microsoft.com/en-us/library/ms191179.aspx

QUESTION NO: 43

You work as a SQL Server 2012 database developer at TestKing.com. TestKing.com has a large database named SalesDB.

New rows are inserted into the tables in the SalesDB database and updates to existing rows occur on a high frequency. The inserts and updates often blocked by queries retrieving and reading data.

How would you prevent queries that retrieve and read data from blocking queries that insert and update data?

A.
You should make use of the SERIALIZABLE ISOLATION LEVEL.

B.

You should make use of the SNAPSHOT ISOLATION LEVEL.

C.
You should make use of the REPEATABLE READ ISOLATION LEVEL.

D.
You should make use of the READCOMMITTED ISOLATION LEVEL.

E.
You should make use of the READPAST ISOLATION LEVEL

Answer: A
Explanation:

Ref: http://msdn.microsoft.com/en-us/library/ms173763.aspx

QUESTION NO: 44

You work as a SQL Server 2012 database developer at TestKing.com. You are developing a stored procedure that updates rows in several tables.

You want the entire transaction to be rolled back should the stored procedure cause a run-time error.

How would you accomplish this?

A.
You should make use of the SET XACT_ABORT ON statement in the stored procedure.

B.
You should have the stored procedure run in the SERIALIZABLE ISOLATION LEVEL.

C.
You should make use of a LOOP hint in the stored procedure.

D.
You should have the stored procedure run in the SNAPSHOT ISOLATION LEVEL.

E.
You should make use of an INSTEAD OF UPDATE trigger in the stored procedure.

F.
You should make use of RAISERROR in the stored procedure.

Answer: A

Explanation:

Ref: http://msdn.microsoft.com/en-us/library/ms188792

QUESTION NO: 45

You work as a database developer at TestKing.com. TestKing has an in-house application named TKApp3 that runs a Transact-SQL query against a SQL Server 2012 database named SalesDB.

TestKing.com users report that TKApp3 is functioning sluggishly. You discover that concurrent updates are causing blockages on the SalesDB database.

How would you ensure that the query to use the original data when updates occur?

A.
You should have the query run in the REPEATABLE READ ISOLATION LEVEL.

B.
You should have the query run in the SERIALIZABLE ISOLATION LEVEL.

C.
You should have the query run in the READCOMMITTED ISOLATION LEVEL.

D.
You should have the query run in the SNAPSHOT ISOLATION LEVEL.

E.
You should have the query run in the READPAST ISOLATION LEVEL.

Answer: D

Explanation:

QUESTION NO: 46

You work as a database developer at TestKing.com. TestKing has an in-house application named TKApp3 that runs a Transact-SQL query against a SQL Server 2012 database.

You want to run an execution plan against the query that will provide detailed information on missing indexes.

How would you accomplish this task?

A.
You should make use of the READPAST hint in the queries.

B.
You should make use of the READCOMMITTED hint in the queries.

C.
You should make use of the SET SHOWPLAN_XML ON statement in the query.

D.
You should make use of the SET STATISTICS XML ON statement in the query.

E.
You should make use of the SET XACT_ABORT OFF statement in the query.

F.
You should make use of the SET CONTEXT_INFO statement in the query.

Answer: C
Explanation:

QUESTION NO: 47

You work as a database administrator at TestKing.com. TestKing.com has a SQL Server 2012 database infrastructure that contains a database named SalesDB.

The SalesDB database is used by an in-house application named TKApp3. TestKing.com users report that TKApp3 is functioning sluggishly.

You discover that application consumes considerable memory when it runs single-use dynamic queries against the SalesDB database. You suspect that these queries are making excessive use of procedure cache.

How would you reduce procedure cache if you cannot create new indexes on the SalesDB database?

A.
You should replace the queries with recursive stored procedures.

B.
You should add make use of the INCLUDE clause in the index.

C.
You should make use of the READPAST hint in the queries.

D.
You should make use of the READCOMMITTED hint in the queries.

E.
You should make use of the optimize for ad hoc workloads option.

Answer: E
Explanation:

QUESTION NO: 48

You work as a database administrator at TestKing.com. TestKing.com has a SQL Server 2012 database infrastructure that contains a database named TestKingDB.

The TestkingDB database is used by an in-house application named TKApp3 that queries a read-only table with a clustered index. TestKing.com users report that TKApp3 is functioning sluggishly.

You suspect query the application uses is causing the problem. You analyze the query and discover that column referenced in the WHERE clause is not part of the clustered index. You also notice that the query returns five columns, as well as a COUNT (*) clause grouped on the five columns.

How would you improve the efficiency of this query?

A.
You should replace the query with recursive stored procedure.

B.
You should replace the COUNT (*) clause with a persisted computed column.

C.
You should create nonclustered indexes on all columns used in the query.

D.
You should create a filtered index on the column used in the WHERE clause.

E.
You should add an INCLUDE clause to the clustered index.

F.
You should create a columnstore index on all columns used in the query.

G.
You should create a unique clustered index on the column used in the WHERE clause.

Answer: F
Explanation:

QUESTION NO: 49

You work as a database administrator at TestKing.com. TestKing.com has a SQL Server 2012 database named ProductsDB. The ProductsDB database is shown in the following database diagram:

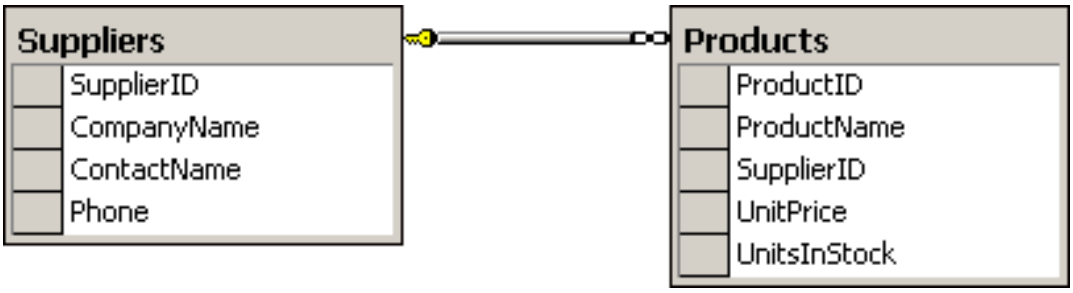

You need to write a Transact-SQL query that displays rows of data in the following XML format:

<Suppliers>

<CompanyName>Company Name</CompanyName>

<ContactName>Contact Name</ContactName>

<Phone>453 3545 5224</Phone>

<Products>

<ProductID>10</ProductID>

<ProductName>Product Name</ProductName>

<UnitPrice>559.00</UnitPrice>

<UnitsInStock>12</UnitsInStock>

</Products>

<Products>

<ProductID>132</ProductID>

<ProductName>Product Name</ProductName>

<UnitPrice>59.00</UnitPrice>

<UnitsInStock>102</UnitsInStock>

</Products>

<Products>

<ProductID>259</ProductID>

<ProductName>Product Name</ProductName>

<UnitPrice>599.00</UnitPrice>

<UnitsInStock>6</UnitsInStock>

</Products>

</Suppliers>

Which of the following SELECT statement would you write?

A.
SELECT s.SupplierID, s.CompanyName, s.ContactName, s.Phone, p.ProductID, p.UnitPrice, p.UnitsInStock

FROM Suppliers AS s

INNER JOIN Products AS p ON s.SupplierID = p.SupplierID

FOR XML AUTO, ELEMENTS

B.
SELECT s.SupplierID, s.CompanyName, s.ContactName, s.Phone, p.ProductID, p.UnitPrice, p.UnitsInStock

FROM Suppliers AS s

INNER JOIN Products AS p ON s.SupplierID = p.SupplierID

FOR XML

C.
SELECT Suppliers.SupplierID, Suppliers.CompanyName, Suppliers.ContactName, Suppliers.Phone, Products.ProductID, Products.UnitPrice, Products.UnitsInStock

FROM Suppliers

INNER JOIN Products ON Suppliers.SupplierID = Products.SupplierID

FOR XML AUTO

D.

```
SELECT Suppliers.SupplierID, Suppliers.CompanyName, Suppliers.ContactName,
Suppliers.Phone, Products.ProductID, Products.UnitPrice, Products.UnitsInStock

FROM Suppliers

INNER JOIN Products ON Suppliers.SupplierID = Products.SupplierID

FOR XML AUTO, RAW
```

Answer: A

Explanation:

QUESTION NO: 50

You work as a SQL Server 2012 database developer at TestKing.com. TestKing.com has a database named SalesDB with tables named Customer and Orders. The Customer and Orders tables were created using the following Transact-SQL code:

```
CREATE TABLE SalesDB.Customers

(

CustomerID int NOT NULL PRIMARY KEY,

CustomerName nvarchar (250) NOT NULL,

Address1 nvarchar (100) NOT NULL,

Address2 nvarchar (100) NULL,

City nvarchar (50) NOT NULL,

State nvarchar (50) NOT NULL,

Zip varchar (5) NOT NULL,

Phone varchar (10) NOT NULL

)

GO

CREATE TABLE SalesDB.Orders

(

OrderID int NOT NULL PRIMARY KEY,
```

CustomerID int NOT NULL,

OrderDate datetime NOT NULL,

ShipDate datetime NOT NULL,

CustomerID int NOT NULL,

SalesRepID int NOT NULL

)

GO

You are developing a stored procedure named OrdersByDate that returns the OrderID, CustomerID, CustomerName and OrderDate. The stored procedure will take a parameter named @date that uses the int datatype. The @date parameter will be used to filter the result set based on the OrderDate column in the Orders table.

How would you create the stored procedure?

A.
CREATE PROCEDURE OrdersByDate

@date int

AS

SELECT OrderID, CustomerID, CustormerName, OrderDate

FROM Orders INNER JOIN Customers ON Orders.CustomerId = Customers.CustomerId

WHERE OrderDate = CONVERT(datetime,@date)

B.
CREATE PROCEDURE OrdersByDate

@date int

AS

SELECT OrderID, CustomerID, CustormerName, OrderDate

FROM Orders INNER JOIN Customers ON Orders.CustomerId = Customers.CustomerId

WHERE OrderDate =@date

C.
CREATE PROCEDURE OrdersByDate

@date int

AS

SELECT OrderID, CustomerID, CustormerName, OrderDate

FROM Orders INNER JOIN Customers ON Orders.CustomerId = Customers.CustomerId

WHERE OrderDate = CAST(@date AS datetime)

D.
CREATE PROCEDURE OrdersByDate

@date int

AS

SELECT OrderID, CustomerID, CustormerName, OrderDate

FROM Orders INNER JOIN Customers ON Orders.CustomerId = Customers.CustomerId

WHERE OrderDate = PARSE(@date AS datetime)

Answer: C
Explanation: